

Amdahl's Law

parallel

concurrent

*thread of control*

# Parallel and concurrent programming

## 2. Programming Model

PROTECTION

data race

synchronization

Michelle Kuttel

divide and conquer algorithms

THREAD SAFETY

correctness

MUTUAL EXCLUSION

locks

liveness

DEADLOCK


starvation

HIGH PERFORMANCE COMPUTING

EXECUTORS

*thread pools*

timing





# Overview: How to write parallel programs

To write a parallel program, you (the programmer) need new **primitives** from a programming language or library, that enable you to:

- run **multiple** operations **at once**
- share **data** between operations
- **coordinate** (*a.k.a. synchronize*) **operations**



# Overview: How to write parallel programs

To write a parallel program, you (the programmer) need new **primitives** from a programming language or library, that enable you to:

- run **multiple** operations **at once**
- share **data** between operations
- **coordinate** (*a.k.a. synchronize*) **operations**

How this works/is done depends on the **parallel programming model** used in the language/library.




Java uses the *Shared Memory*  
parallel programming model



# The Shared Memory Model

All memory is placed into a single (physical) address space.


- Processors connected by some form of **interconnection network**
- Single virtual address space across all of memory.** Each processor can access all locations in memory.





# Shared Memory


The **ideal** picture of shared memory:





# Shared Memory


The **ideal picture** of shared memory:






# Shared Memory

The **ideal picture** of shared memory:




The **actual architecture** of shared memory systems:

*Symmetric Multi-Processor (SMP):*



*Distributed Shared Memory (DSM):*




[Also have Non-uniform Memory Access Symmetric Multi-Processor (NUMA-SMP)]





## An aside: cache (Architecture)



A memory cache, also called a "CPU cache," is a memory bank that bridges main memory and the processor.

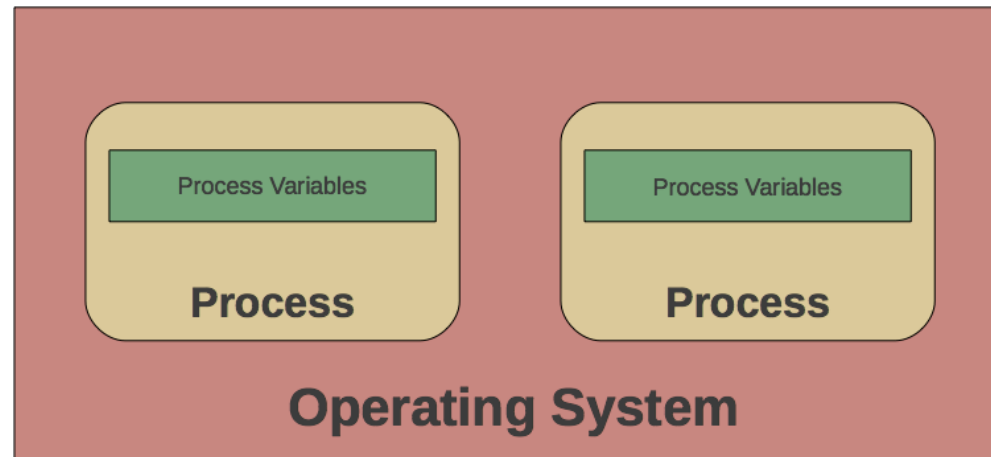
- It has faster static RAM (SRAM) chips than the dynamic RAM (DRAM) used for main memory.
- The cache allows instructions to be executed and data to be read and written at **higher speed**.

### Can have multiple caches (L1, L2, L3) in modern chips

- L1 is the fastest; each subsequent cache is slower and larger than L1, and instructions and data are staged from main memory to L3 to L2 to L1 to the processor.
- On multicore chips, the L3 cache is generally shared among all the processing cores.




## A Process (operating system)



Operating system unit of resource allocation both for CPU time and for memory.

- A process is represented by its **code**, **data** and the **state of the machine registers**.
- data of the process divided into
  - global variables and local variables
  - organized as a **stack**.
- Generally, each process in an operating system has its own address space
  - entirely separate entities

It is hard to obtain parallelism or  
concurrency with separate  
*processes...* (why?)



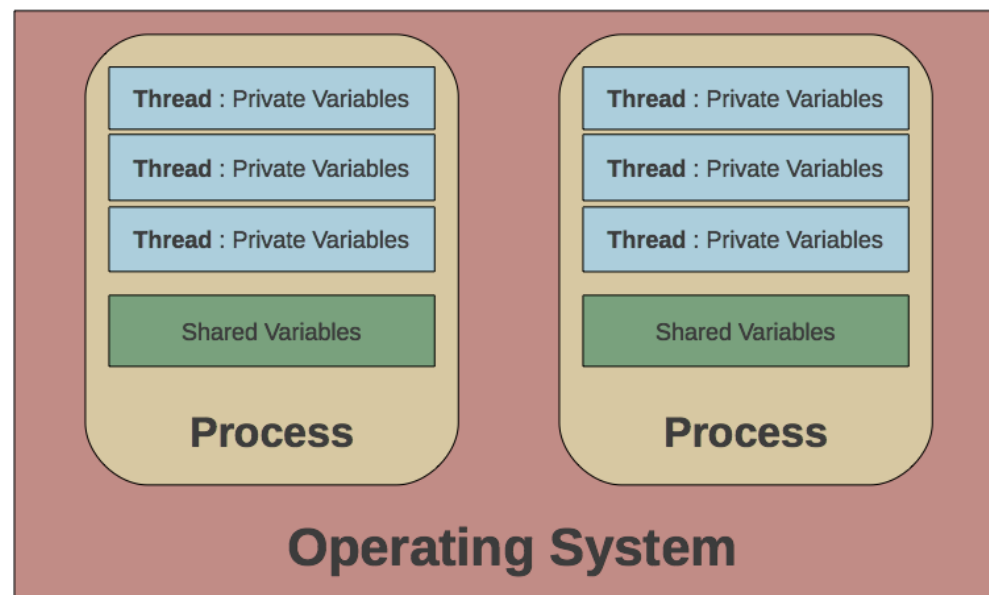
... so operating systems created  
*threads.*

# Thread

Process given internal concurrency with multiple *lightweight processes* or ***threads***.

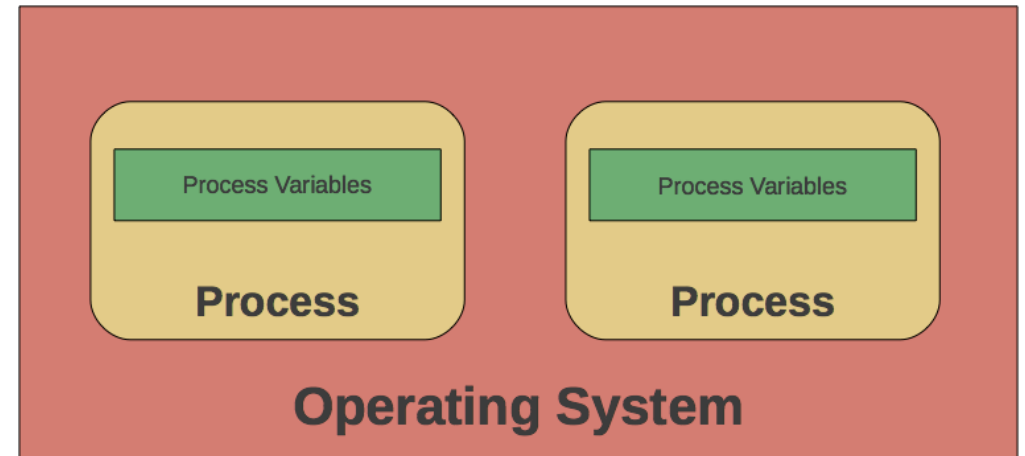
- multiple ***threads of control***
- a process with multiple (lightweight) threads of control has multiple stacks
  - one for each thread.

but access to **shared memory** too.

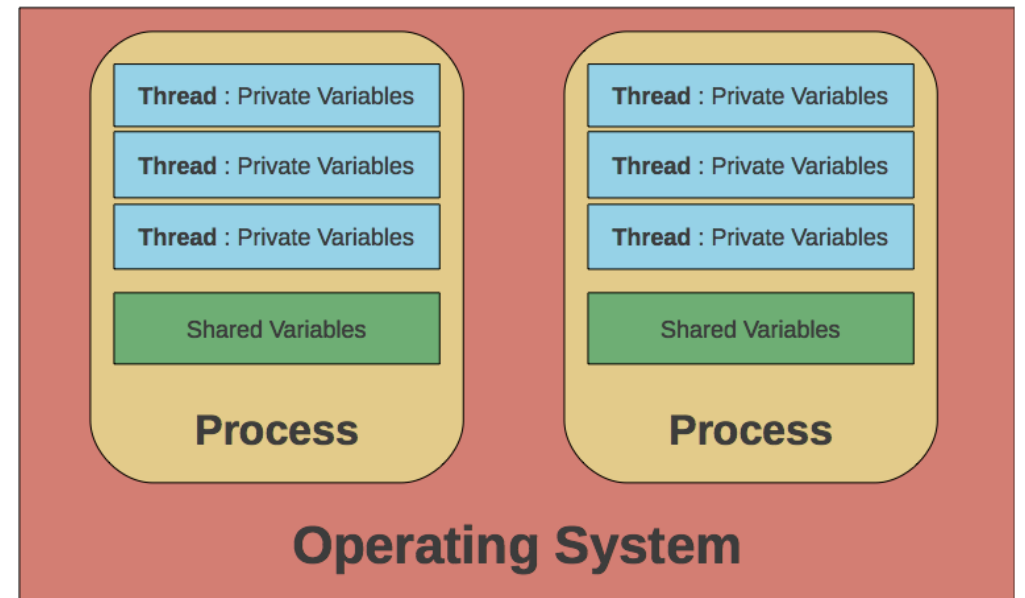


# Processes versus threads

## Process memory model



## Thread memory model



# What is a parallel program?

The **shared memory model** has multiple **explicit threads** running concurrently.

Threads can:

- perform **multiple computations** in parallel;
- perform separate **simultaneous activities**;
- **communicate easily** and **implicitly** with each other through **shared memory**.

(but this is **dangerous** if you don't protect your variables correctly)

\* This is true for the shared memory model of parallel computing discussed in this module.

# Programming Model: Sequential program state

A running serial program has

- One *program counter* (current statement executing)
- One *call stack*
  - each stack frame holds the local variables for a method call that has started but not yet finished.
  - Calling a method pushes a new frame and returning from a method pops a frame.
  - Call stacks are why recursion is not “magic.”
- *Objects*. Object are created by calling **new**. We call the memory that holds all the objects the *heap*.  
(nothing to do with data structure called a heap)
- *Static fields* of classes.




# Programming Model: Shared memory


- Each thread has its own program counter, call stack and local variables
- All threads share one collection of objects and static fields
- Static fields of classes are also shared by all threads.
- Threads communicate through shared objects (implicit communication)
  - To *communicate*, **write** somewhere another thread **reads**



# Sequential Computation



# Concurrent Computation






# How threads run

As the programmer, **you create threads.**

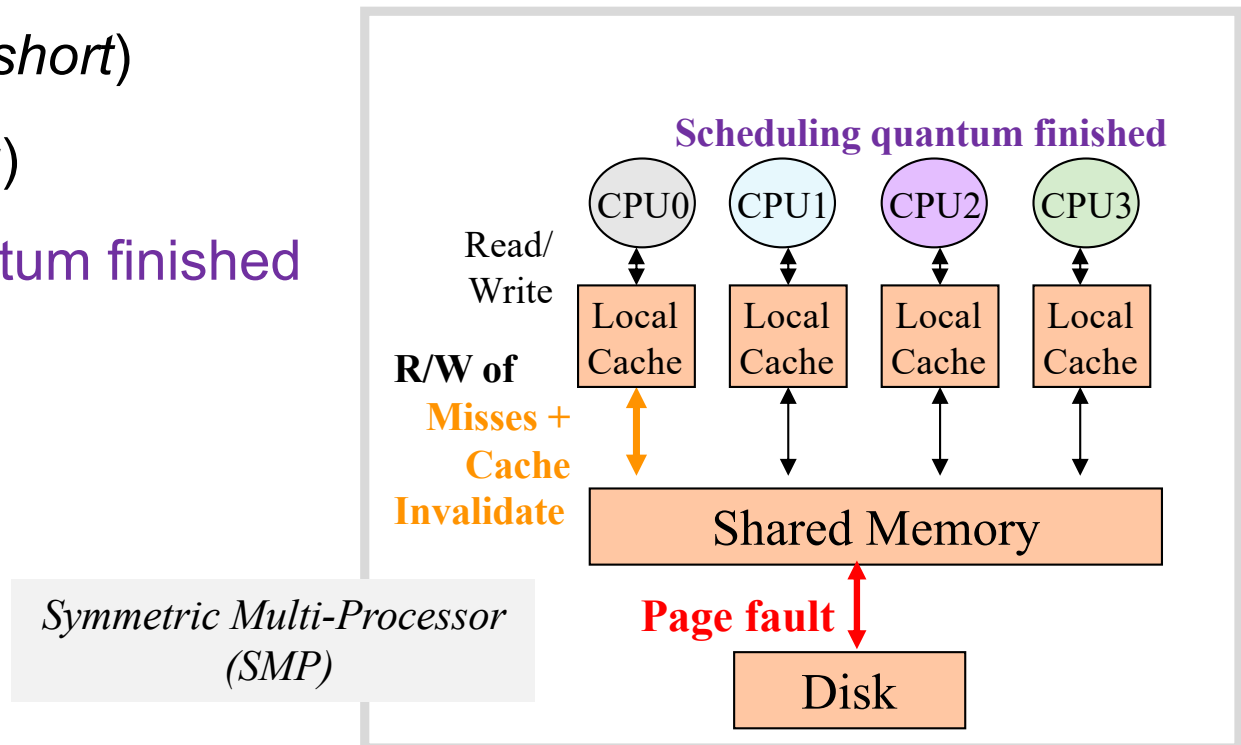
- The **Operating System scheduler** determines how and when those threads are run on the available processors
  - Unless you are writing a scheduler, **you don't have control over this**
  - You don't know how many **processors/cores** your program will use  
(though you can guess).
  - You don't know **the order** in which threads will execute  
(you can't even guess).

# Asynchrony



Threads are subject to **sudden unpredictable delays**:

- **Cache misses** (*short*)
- **Page faults** (*long*)
- **Scheduling quantum finished** (*really long*)




# Other parallel programming models

We focus on **shared memory**, but several **other models** exist. Common alternatives are:

- **Message-passing:**

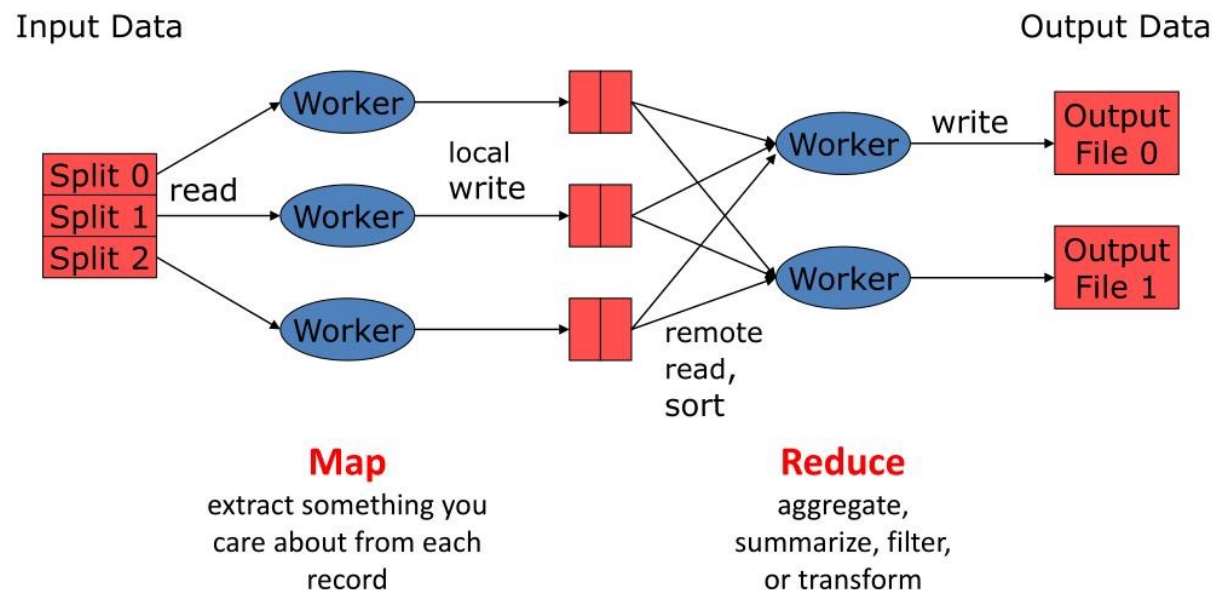
- Explicit **threads/processes**, each with their own objects/data.
- Communication is via explicitly **sending/receiving messages**, containing **copies** of the data (share nothing)
- Most common model on HPC systems (though usually in a hybrid with shared memory)



# Other parallel programming models

- **Map reduce:**

- Data parallelism concept from functional programming languages like LISP
- Have **primitives** for things like “apply function to every element of an array in parallel”.
  - details of the **underlying parallelization** are **hidden** from the programmer, provided you can express your program using the available primitives
- MapReduce was developed by Google and the programming model has since been adopted by many software frameworks, e.g. Apache’s open-source Hadoop





# Other parallel programming models

## Golang

- Go-routines,
  - Go runtime maps these onto operating system threads
- Channels – used for communication between Go-routines

```
package main
```

```
import "fmt"
```

```
func main() {  
    fmt.Println("Hello, World!")  
}
```