

parallel
concurrent
Amdahl's Law
thread of control
processors versus processes

Parallel and concurrent programming

7.

protection data race synchronization
divide-and-conquer algorithms

Part II: Concurrency and Mutual Exclusion

THREAD
correctness
SAFETY

Michelle Kuttel

MUTUAL EXCLUSION

locks

readers-writers problem

liveness

DEADLOCK

starvation

HIGH PERFORMANCE COMPUTING

EXECUTORS


thread pools

producer-consumer problem

timing

Dining philosophers problem

Multithreaded programming




Parallel programming

So far focused on **parallel algorithms**

- used fork-join only
 - Lower *span* via parallel tasks
 - threads all do the same thing, really, on different data
 - “**data decomposition**”


Algorithms all had a very *simple structure* to avoid race conditions

- Each thread had memory “only it accessed”
 - Example: array sub-range
- On **fork**, allocated some of its memory to subthread and did not access that memory again until after **join** on the subthread
 - Only form of synchronization used, but vital for fork-join library.





Now, concurrent programming



Threads assigned *independent* tasks

- access same resources (rather than implementing the same algorithm)
- “**task decomposition**”


Concurrency: Correctly and efficiently managing access to shared resources from multiple (possibly simultaneous) clients

More **complex structure** for algorithms

- multiple threading doing different things with data
- e.g. animator threads, threads that update data.



Now, concurrent programming



Threads assigned *independent* tasks

- access same resources (rather than implementing the same algorithm)
- “**task decomposition**”

Concurrency: Correctly and efficiently managing access to shared resources from multiple (possibly simultaneous) clients

More **complex structure** for algorithms

- multiple threading doing different things with data
- e.g. animator threads, threads that update data.

Race conditions much more prevalent!



Concurrent programming

Unlike parallelism, not all about implementing algorithms faster. :

- ***Different tasks may need to happen at once for program to work.***
 - Example: Animation in a game
- ***Simultaneous tasks may be necessary for responsiveness***
 - Example: Respond to GUI events in one thread while another thread is performing an expensive computation
 - Web server
- ***Simultaneous tasks may be necessary for efficient processor utilization***
 - If 1 thread “goes to disk,” have something else to do (often called “latency hiding”)
- ***Simultaneous tasks may be necessary for failure isolation***
 - Convenient structure if want to *interleave* multiple tasks and don’t want an exception in one to stop the other

Synchronization

Synchronization constraints are **requirements** pertaining to the **order of events**.

This is the programmer's responsibility to ensure that synchronization constraints are enforced

- “The compiler” has no idea what interleavings should or shouldn't be allowed in your program
- you need language primitives to do it!

Simple example:

- **Signaling:**

Event A must happen **before** **Event B**.

- We've seen this already – solved with `join`.
 - Not a universal solution – why?

Synchronization

- In **real life** we typically use a **clock** to enforce synchronization constraints.
 - How do we know if A happened before B?
 - If we know what time both events occurred, we can just compare the times.
- In computer systems, we need to satisfy synchronization constraints without a clock
 - there is **no universal clock across different computers**
 - we don't know with fine enough **resolution** when events occur.
 - (clock ticks are too far apart)



A Mutual Exclusion Fable: Alice and Bob share a garden... and have pets



Mutual
exclusion:

Event A
and Event B
must **not**
happen ***at
the same
time.***



The Mutual Exclusion problem

One of the most important synchronization problems in concurrent programming

- several processes compete for a resource, but the nature of the resource **requires** that only one process at a time actually accesses the resource at any point
- use **synchronization** to avoid **incorrect simultaneous access**

Also called the **critical section problem**.

Abstraction of many synchronization issues.

If two processes try to enter a **critical section**, one process must win and the other must **block**:

- not proceed until the “winning” process has completed



Examples of where mutual exclusion is necessary

Multiple threads:

1. Processing different bank-account operations
 - What if 2 threads change the same account at the same time?
2. Using a shared cache (e.g., hash table) of recent files
 - What if 2 threads insert the same file at the same time?
3. Creating a pipeline (like an assembly line) with a queue for handing work to next thread in sequence
 - What if enqueueer and dequeuer adjust a circular array queue at the same time?

Simple example: counter shared between threads

```
public class Counter {  
    private long value;  
  
    Counter() { value=0; }  
  
    public long get() {return value;}  
  
    public void set(long newVal) {  
        value=newVal;  
    }  
  
    public void incr() {value++;}  
}
```

Procedure for Thread i

```
public class CounterUpdateThread extends Thread {
```

```
    Counter sharedCount;
```

Shared counter
object

```
    int rep;
```

```
    CounterUpdateThread(Counter c, int repeats) {
```

```
        sharedCount= c;
```

```
        rep=repeats;
```

```
    }
```

```
    public void run() {
```

```
        for (int i=0; i<rep; i++) {
```

```
            sharedCount.incr();
```

```
        }
```

```
    }
```

```
}
```

Testing it all

```
public class TestCounterSafety {
    public static void main(String args[]) throws InterruptedException
    {
        int noThrds = 100;
        int addPerThread=100;
        Counter sharedCount= new Counter();

        CounterUpdateThread [] thrds= new CounterUpdateThread[noThrds];
        for (int i=0;i<noThrds;i++) {
            thrds[i]=new CounterUpdateThread(sharedCount,addPerThread);
        }
        for (int i=0;i<noThrds;i++) {
            thrds[i].start();
        }

        for (int i=0;i<noThrds;i++) {
            thrds[i].join();
        }

        int expectedVal = noThrds*addPerThread;
        System.out.println("Final value of counter is:"
            + sharedCount.get() + " and should be:" + expectedVal);
    }
}
```

Testing it all

```
public class TestCounterSafety {
    public static void main(String args[]) throws InterruptedException
    {
        int noThrds = 100;
        int addPerThread=100;
        Counter sharedCount= new Counter();

        CounterUpdateThread [] thrds= new CounterUpdateThread[noThrds];
        for (int i=0;i<noThrds;i++) {
            thrds[i]=new CounterUpdateThread(sharedCount,addPerThread);
        }
        for (int i=0;i<noThrds;i++) {
            thrds[i].start();
        }

        for (int i=0;i<noThrds;i++) {
            thrds[i].join();
        }

        int expectedVal = noThrds*addPerThread;
        System.out.println("Final value of counter is:"
            + sharedCount.get() + " and should be:" + expectedVal);
    }
}
```

Race condition!

Simple example: counter shared between threads

```
public class Counter {  
    private long value;  
  
    Counter() { value=0; }  
  
    public long get() {return value;}  
  
    public void set(long newVal) {  
        value=newVal;  
    }  
  
    public void incr() {value++;}  
}
```

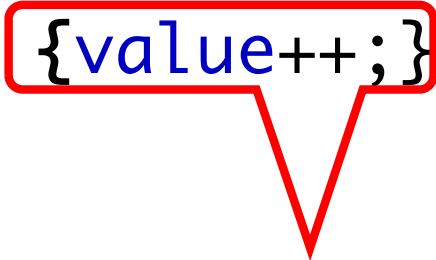
OK for single
thread,
not for concurrent
threads

What It Means

```
public class Counter {  
    private long value;  
    //[...]  
    public void incr() {value++;}  
}
```

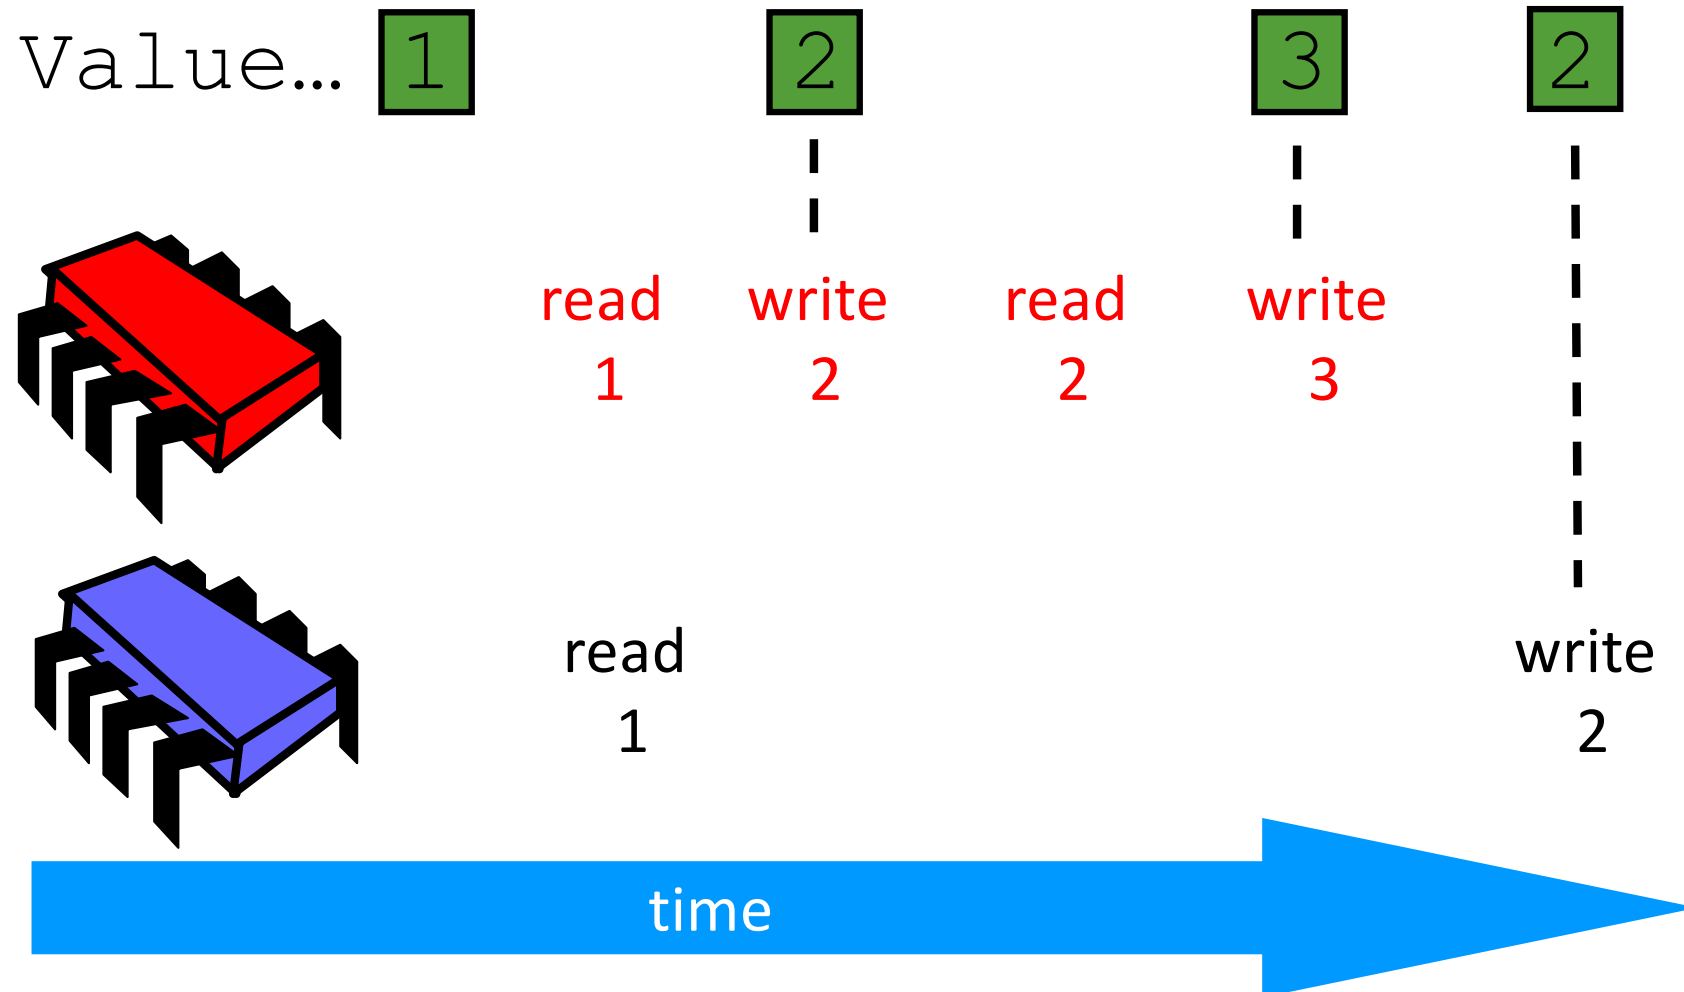
What It Means

```
public class Counter {  
    private long value;  
    // [...]   
    public void incr() {value++;}  
}
```



```
temp = value;  
temp = temp + 1;  
return temp;
```

Not so good...





Solving Mutual exclusion in Java

`join` is not what we want

- We need to block until another thread is “done using what we need” not “completely done executing”



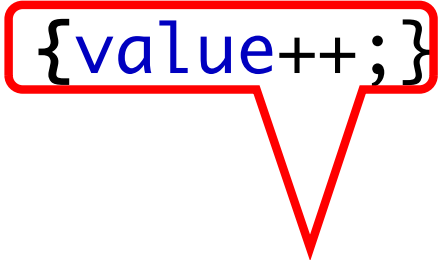
We need atomic actions

Operations A and B are **atomic** with respect to each other if, from the perspective of a thread executing A, when another thread executes B, either all of B has executed or none of it has.

The operation is *indivisible*.

What It Means

```
public class Counter {  
    private long value;  
    //[...]  
    public void incr() {value++;}  
}
```



Make these steps
atomic (indivisible)

```
temp = value;  
temp = temp + 1;  
return temp;
```

How do we do this?

Mutual exclusion: Atomic classes in Java

One way we can fix this is with a thread-safe **atomic variable class**: `java.util.concurrent` contains **atomic variable classes** for atomic actions on numbers and objects.

- `AtomicInteger`, `AtomicLong`, `AtomicBoolean`, `AtomicReference`
- **note: Integer NOT atomic by default.**
- **no atomic classes** for character and double.
- all have `getAndSet()` method that atomically sets the value and returns it

So, we can replace a long counter with `AtomicLong` to ensure that all updates to the counter state are **atomic**.

CORRECT Procedure for Thread *i*

```
public class Counter {  
    private AtomicLong value; // this is a class!  
  
    Counter() { value=new AtomicLong(0); }  
  
    public long get() { return value.get(); }  
  
    public void set(long val) { value.set(val); }  
  
    public void incr() { value.getAndIncrement(); }  
  
}
```

Seems to work...!



Limitations of Atomic variables


Atomic variables only make a class thread-safe only if ONE variable defines the class **state** and there are no **compound actions** accessing this **state**.

- to preserve **state consistency**, you should update related state variables in a **single** atomic operation

Example 2: atomic won't work

```
class BankAccount {
    private int balance = 0;
    int  getBalance() { return balance; }
    void setBalance(int x) { balance = x; }

    // withdraw is a compound action
    void withdraw(int amount) {
        int b = getBalance();
        if (amount > b)
            throw new WithdrawTooLargeException();
        setBalance(b - amount);
    }
    ... // other operations like deposit, etc.
}
```



Interleaving

Suppose:

- Thread **T1** calls **`x.withdraw(100)`**
- Thread **T2** calls **`y.withdraw(100)`**

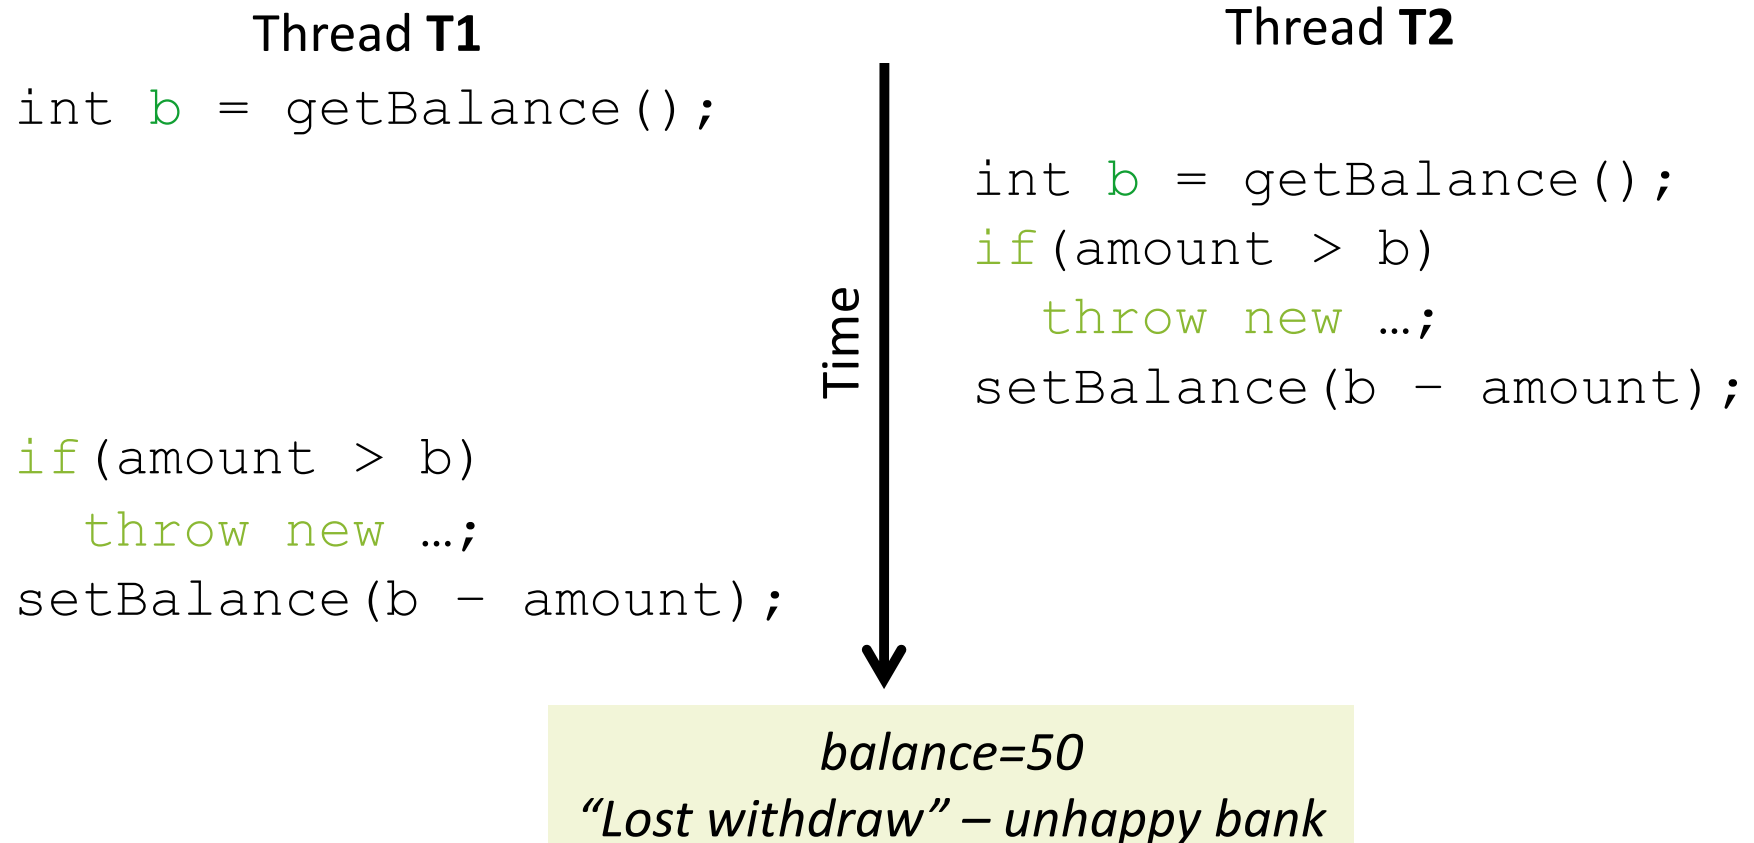
If calls **interleave**

- and **`x`** and **`y`** refer to different accounts...
 - > no data race
- and **`x`** and **`y`** aliases...
 - > possible data race

Race condition: A bad interleaving

Interleaved `withdraw(100)` calls on the same account

- Assume initial `balance == 150`



Incorrect “fix”

It is tempting and **almost always wrong** to fix a bad interleaving by rearranging or repeating operations, such as:

```
void withdraw(int amount) {  
    if (amount > getBalance())  
        throw new WithdrawTooLargeException();  
    // maybe balance changed  
    setBalance(getBalance() - amount);  
}
```

Incorrect “fix”

```
void withdraw(int amount) {  
    int b = getBalance();  
    if(amount > b)  
        throw new WithdrawTooLargeException();  
    setBalance(b - amount);  
}
```

It is tempting and **almost always wrong** to fix a bad interleaving by rearranging or repeating operations, such as:

```
void withdraw(int amount) {  
    if(amount > getBalance())  
        throw new WithdrawTooLargeException();  
    // maybe balance changed  
    setBalance(getBalance() - amount);  
}
```

This fixes nothing!

- Narrows the problem by one statement
- (Not even that, since the compiler could turn it back into the old version because you didn't indicate need to synchronize)
- And now a negative balance is possible – why?



Mutual exclusion is required

- At most one thread withdraws from account **A** at a time
- Exclude other simultaneous operations on **A** too (e.g., deposit)

Mutual exclusion in Java: synchronized methods

A built-in locking mechanism in Java for enforcing **atomicity**.

A method/ function can be made **mutually exclusive** by prefixing the method with the keyword **synchronized**.

```
synchronized void f() { body; }
```

Java implementation

```
class BankAccount {  
    private int balance = 0;  
    synchronized int getBalance()  
    { return balance; }  
    synchronized void setBalance(int x)  
    { balance = x; }  
    synchronized void withdraw(int amount) {  
        int b = getBalance();  
        if(amount > b)  
            throw ...  
        setBalance(b - amount);  
    }  
    // deposit would also use synchronized  
}
```

Can also be used in example 1:
Safe Counter Implementation v2

```
public class Counter {  
    private long value;  
    //[...]  
    public synchronized void incr() {value++;}  
}
```

But atomic class should be faster
for this simple example... as
avoids actual synchronization
behind the scenes and can use
hardware support.

Mutual exclusion in Java: synchronised block

Critical sections of code can be made **mutually exclusive** by prefixing the method with the keyword **synchronized**.

- `synchronized void f() { body; }` is equivalent to:
- `void f() { synchronized(this) { body; } }`
- a **synchronized block** has two parts:
 - a reference to an object that will serve as the *lock*
 - a block of code to be guarded by that *lock*

```
synchronized (object) {  
    statements }
```




How does this work? - Locks

Every object (but not any primitive type) “is a lock” in Java

- A lock is an Abstract Data Type with operations:
 - **new**: make a new lock, initially “*not held*”
 - **acquire**: blocks if this lock is already currently “*held*”
 - Once “*not held*”, makes lock “*held*”
 - **release**: makes this lock “*not held*”
 - if ≥ 1 threads are blocked on it, exactly 1 will acquire it
- The lock implementation ensures that upon **simultaneous acquires and/or releases**, a correct thing will happen
 - Example: Two acquires: one will “win” and one will block
- How can this be implemented?
 - Need to “check and update” “all-at-once”
 - Uses special hardware and O/S support
 - Covered in operating systems module in CSC3002F

You can use Java's locks directly

```
class BankAccount {  
    private int balance = 0;  
    private Lock lk = new Lock();  
    ...  
    void withdraw(int amount) {  
        lk.acquire(); /* may block */  
        int b = getBalance();  
        if(amount > b)  
            throw new WithdrawTooLargeException();  
        setBalance(b - amount);  
        lk.release();  
    }  
    // deposit would also acquire/release lk  
}
```

Problems with locks

- A lock is a very **primitive mechanism**
 - Still up to you to use correctly to implement critical sections
- **Incorrect:** Use different locks for **withdraw** and **deposit**
 - Mutual exclusion works only when using **same** lock
- Poor performance: Use same lock for every bank account
 - No simultaneous operations on different accounts
- Incorrect: Forget to release a lock (blocks other threads forever!)
 - Previous slide is **wrong** because of the exception possibility!

```
if (amount > b) {  
    lk.release(); // hard to remember!  
    throw new WithdrawTooLargeException();  
}
```

Other issues with locks

- If **withdraw** and **deposit** use the same lock, then simultaneous calls to these methods are properly synchronized
- But what about **getBalance** and **setBalance**?
 - Assume they're **public**, which may be reasonable
- If they don't acquire the same lock, then a race between **setBalance** and **withdraw** could produce a wrong result
- If they do acquire the same lock, then **withdraw** would block forever because it tries to acquire a lock it already has

```
class BankAccount {  
    private int balance = 0;  
    synchronized int getBalance()  
    { return balance; }  
    synchronized void setBalance(int x)  
    { balance = x; }  
    synchronized void withdraw(int amo  
        int b = getBalance();  
        if(amount > b)  
            throw ...  
        setBalance(b - amount);  
    }  
    // deposit would also use synchroniz
```

Locks have to be *re-entrant*

```
int setBalance(int x) {  
    lk.acquire();  
    balance = x;  
    lk.release();  
}  
  
void withdraw(int amount) {  
    lk.acquire();  
    ...  
    setBalance1(b - amount);  
    lk.release();  
}
```

This simple code works fine provided **lk** is a reentrant lock

- Okay to call **setBalance** directly
- Okay to call **withdraw** (won't block forever)

Re-entrant lock

A **re-entrant** lock (a.k.a. **recursive lock**)

- “Remembers”
 - the thread (if any) that currently holds it
 - a *count*
- When the lock goes from *not-held* to *held*, the count is 1
- If (code running in) the current holder calls **acquire**:
 - it does not block
 - it increments the count
- On **release**:
 - if the count is > 0 , the count is decremented
 - if the count is 0, the lock becomes *not-held*

Reentrant locks

- This code would deadlock without the use of reentrant locks:

```
class Widget {  
    public synchronized void doSomething() {  
        ....  
    }  
}  
  
class BobsWidget {  
    public synchronized void doSomething() {  
        System.out.println("Calling super");  
        super.doSomething();  
    }  
}
```

Java does it for you: synchronization

Every Java object created, including every Class loaded, has an associated **lock** (or **monitor**).

- Putting code inside a synchronized block makes the compiler append instructions to **acquire** the lock on the specified object before executing the code, and **release** it afterwards (either because the code finishes normally or abnormally).
- Between acquiring the lock and releasing it, a thread is said to "own" the lock. **At the point of Thread A** wanting to **acquire the lock**, if **Thread B** already **owns it**, then **Thread A must wait for Thread B to release it.**

Java locks summary

Every Java object possesses **one lock**

- Manipulated only via `synchronized` keyword
- **Class objects** contain a lock used to protect statics
- Scalars like `int` are not Objects, so can only be locked via their enclosing objects

Java locks are **reentrant**

- A thread hitting `synchronized` passes if the lock is free or it already possesses the lock, else waits
- Released after passing as many `}`'s as `{`'s for the lock
 - cannot forget to release lock



Who gets the lock?

- There are **no fairness specifications in Java**, so if there is contention to call synchronized methods of an object, an arbitrary process will obtain the lock.



Java: block synchronization versus method synchronization

Block synchronization is preferred over **method synchronization**:

- with block synchronization you only need to lock the critical section of code, instead of whole method.
- Synchronization comes with a performance cost:
 - we should synchronize only part of code which absolutely needs to be synchronized.

Why we need locks

We can't implement our own mutual-exclusion protocol.

```
class BankAccount {
    private int balance = 0;
    private boolean busy = false;
    void withdraw(int amount) {
        while(busy) { /* "spin-wait" */ }
        busy = true;
        int b = getBalance();
        if(amount > b)
            throw new WithdrawTooLargeException();
        setBalance(b - amount);
        busy = false;
    }
    // deposit would spin on same boolean
}
```

Still just moved the problem!

Thread 1

```
while(busy) { }  
  
busy = true;  
  
int b = getBalance();  
  
if(amount > b)  
    throw new ...;  
setBalance(b - amount);
```

Thread 2

```
while(busy) { }  
  
busy = true;  
  
int b = getBalance();  
if(amount > b)  
    throw new ...;  
setBalance(b - amount);
```

**“Lost withdraw” –
unhappy bank**